



**1. A plot shows episode rewards of at least 800 training episodes in LunarLander-v2 (5%)**



**2. A plot shows episode rewards of at least 800 training episodes in LunarLanderContinuous-v2 (5%)**





Update behavior network 是從replay memory中隨機sampling遊戲過程中儲存的 Next state, current state, action, reward,  $(\phi_t, a_t, r_t, \phi_{t+1})$  這些資訊來做TD-Learning, 再計算 Q value與Q target的MSE loss

$$Q_{target} = \gamma_j + \max \hat{Q}(\phi_{j+1}, a'; \theta^-)$$

Perform a gradient descent step on  $(y_j - \hat{Q}(\phi_{j+1}, a'; \theta^-))^2$  with respect to the network parameter  $\theta$

```
class DQN:
    def _update_behavior_network(self, gamma):
        # sample a minibatch of transitions
        state, action, reward, next_state, done =
            self._memory.sample(self.batch_size, self.device)

        ## TODO ##
        q_value = self._behavior_net(state).gather(dim=1, index=action.long())
        with torch.no_grad():
            q_next = self._target_net(next_state).max(dim=1)[0].view(-1, 1)
            q_target = reward + gamma*q_next*(1-done)
        criterion = nn.MSELoss()
        loss = criterion(q_value, q_target)

        # bp
        self._optimizer.zero_grad()
        loss.backward()
        nn.utils.clip_grad_norm_(self._behavior_net.parameters(), 5)
        self._optimizer.step()
```

每隔一段時間, 就用behavior network 取代 target network

```
class DQN:
    def _update_target_network(self):
        '''update target network by copying from behavior network'''
        ## TODO ##
        self._target_net.load_state_dict(self._behavior_net.state_dict())
```

## Deep Deterministic Policy Gradient (DDPG):

Deep Deterministic Policy Gradient是一個model-free, off-policy的演算法, 他跟DQN的差別有兩個, 第一個是可以解決連續動作的問題, 因為DQN的output是action dim, 是離散的, DDPG的output是連續的, 我是取 $\tanh(x)$ (如下圖), 第二個是DQN是value based的方法去估計 $Q(s,a)$ , 但DDPG是actor-critic的方法, 所以有actor network跟critic network(如下圖)

Actor network的目的是要產生連續動作action, 範圍是  $-1 \sim +1$ , 所以取 $\tanh$

Critic network的目的是要評估在這個state下, actor所採取的action好不好, 所以他跟DQN很像, 都是根據state action pair去估計value的

```
class ActorNet(nn.Module):
    def __init__(self, state_dim=8, action_dim=2, hidden_dim=(400, 300)):
        super().__init__()
        ## TODO ##
        self.fc1=nn.Linear(state_dim,hidden_dim[0])
        self.fc2=nn.Linear(hidden_dim[0],hidden_dim[1])
        self.fc3=nn.Linear(hidden_dim[1],action_dim)
        self.relu=nn.ReLU()
        self.tanh=nn.Tanh()

    def forward(self, x):
        ## TODO ##
        out=self.relu(self.fc1(x))
        out=self.relu(self.fc2(out))
        out=self.tanh(self.fc3(out))
        return out

class CriticNet(nn.Module):
    def __init__(self, state_dim=8, action_dim=2, hidden_dim=(400, 300)):
        super().__init__()
        h1, h2 = hidden_dim
        self.critic_head = nn.Sequential(
            nn.Linear(state_dim + action_dim, h1),
            nn.ReLU(),
        )
        self.critic = nn.Sequential(
            nn.Linear(h1, h2),
            nn.ReLU(),
            nn.Linear(h2, 1),
        )

    def forward(self, x, action):
        x = self.critic_head(torch.cat([x, action], dim=1))
        return self.critic(x)
```

在select action的時候, action network選擇action並加上noise, 增加隨機性, 因為有時候從既有的state去找是很難找到best action的

```
#DDPG
def select_action(self, state, noise=True):
    '''based on the behavior (actor) network and exploration noise'''
    ## TODO ##
    with torch.no_grad():
        if noise:
            re = self._actor_net(torch.from_numpy(state).view(1,-1).to(self.device))+\
                torch.from_numpy(self._action_noise.sample()).view(1,-1).to(self.device)
        else:
            re = self._actor_net(torch.from_numpy(state).view(1,-1).to(self.device))
    return re.cpu().numpy().squeeze()
```

在training過程中, 要更新behavior的actor network  $\mu$ , critic network  $Q$ , target的actor network  $\mu'$ , critic network  $Q'$ , 利用target network output的Q\_target 與behavior Network output 的 Q value做MSE loss更新Q function

```
## update critic ##
# critic loss
## TODO ##
q_value = self._critic_net(state,action)
with torch.no_grad():
    a_next = self._target_actor_net(next_state)
    q_next = self._target_critic_net(next_state,a_next)
    q_target = reward + gamma*q_next*(1-done)
criterion = nn.MSELoss()
critic_loss = criterion(q_value, q_target)
```

behavior Network 的 critic network output 的  $Q(s,a)$ , 我們希望 $Q(s,a)$ 越大越好,  $Q(s,a)$ 越大代表在這個state下越適合採取這個action, 因此定義

$$Actor\ Loss = E[-Q(s, \mu(s))]$$

```
## update actor ##
# actor loss
## TODO ##
action = self._actor_net(state)
actor_loss = -self._critic_net(state,action).mean()
```

#### 4. Describe differences between your implementation and algorithms.(10%)

在training的時候, 每一個episode都會有一段warmup時間(10000 steps),在這段時間中不會去update network得參數, 只會隨便探索, 並把探索的 state, next state, reward 存到replay memory裡。

#### 5. Describe your implementation and the gradient of actor updating.(10%)

Actor network 會輸出一個action, critic network會根據這個action輸出一個的  $Q(s,a)$ , 為了使  $Q(s,a)$  value最大, 因此定義actor loss

$$Actor\ Loss = -Q(s, \mu(s))$$

$$\frac{\nabla L}{\nabla \theta_u} = -\frac{\nabla Q(s, a|\theta_u)}{\nabla a} \times \frac{\nabla a}{\nabla u(s|\theta_u)} \times \frac{\nabla u(s|\theta_u)}{\nabla \theta_u} = -\frac{\nabla Q(s, a|\theta_u)}{\nabla u(s|\theta_u)} \times \frac{\nabla u(s|\theta_u)}{\nabla \theta_u}$$

```
## update actor ##
# actor loss
## TODO ##
action = self._actor_net(state)
actor_loss = -self._critic_net(state,action).mean()
```

#### 6. Describe your implementation and the gradient of critic updating.(10%)

利用target network output 的  $Q$  target 與 behavior network output 的  $Q(s,a)$ 做 mean square error來更新  $Q$  function

$$Critic\ Loss = \frac{1}{N} \sum (Q_{target} - Q(s_t, a_t|\theta_Q))^2$$

```
## update critic ##
# critic loss
## TODO ##
q_value = self._critic_net(state,action)
with torch.no_grad():
    a_next = self._target_actor_net(next_state)
    q_next = self._target_critic_net(next_state,a_next)
    q_target = reward + gamma*q_next*(1-done)
criterion = nn.MSELoss()
critic_loss = criterion(q_value, q_target)
```

#### 7. Explain effects of the discount factor.(5%)

在model free中, 會用未來會得到的reward對現在進行估計, 因此希望越接近現在reward影響越大, 離現在越遠的reward影響比較小, 所以加入discount factor  $\lambda$

$$G_t = R_{t+1} + \lambda R_{t+2} + \lambda^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \lambda^k R_{t+k+1}$$

## 8. Explain benefits of epsilon-greedy in comparison to greedy action selection.(5%)

為了確保在explore和exploit之間取得平衡, 因此使用epsilon-greedy, 有 $\epsilon$ 的機率會去explore未知的action, 有 $1 - \epsilon$ 的機率會選擇現在最佳的action, 因為如果一直選擇現在最佳的action則有可能找不到比這個更好的action

## 9. Explain the necessity of the target network.(5%)

Target Network的目的是為了使輸出更穩定, 所以每隔一段時間, 當Q target改變的時候, 就把behavior network 複製到target network, 使輸出更穩定

## 10. Explain the effect of replay buffer size in case of too large or too small.(5%)

每一個episode都會有一段warmup時間(10000 steps), warmup會儲存到replay buffer, 如果replay buffer size越大, training過程會越穩定, 但會降低training的速度, 如果replay buffer size越小, 容易造成overfitting, 很不穩定, 容易train失敗

## 11. Bonus - Implement and experiment on Double-DQN (10%)

Double-DQN 跟DQN的差別在於, 在update behavior network是如何決定Q target的, DQN的缺點是估計Q target時會過度估計, Double-DQN在估計Q target不是直接用 $\max Q'(s,a)$ , 而是用 $Q(s,a_i)$ 中最大值的i作為t查找 $Q'(s,a_i)$  index

DQN:

$$y_j = R_j + \gamma \max Q'(\phi(S'_j), A'_j, w')$$

Double-DQN:

$$y_j = R_j + \gamma \max Q'(\phi(S'_j), \operatorname{argmax} Q(\phi(S'_j), a, w), w')$$

```
class DDQN:
    def _update_behavior_network(self, gamma):
        # sample a minibatch of transitions
        state, action, reward, next_state, done =
            self._memory.sample(self.batch_size, self.device)

        ## TODO ##
        q_value = self._behavior_net(state).gather(dim=1, index=action.long())
        with torch.no_grad():
            action_index = self._behavior_net(next_state).max(dim=1)[1].view(-1, 1)
            q_next = self._target_net(next_state).gather(dim=1, index=action_index.long())
            q_target = reward + gamma*q_next*(1-done)
        criterion = nn.MSELoss()
        loss = criterion(q_value, q_target)

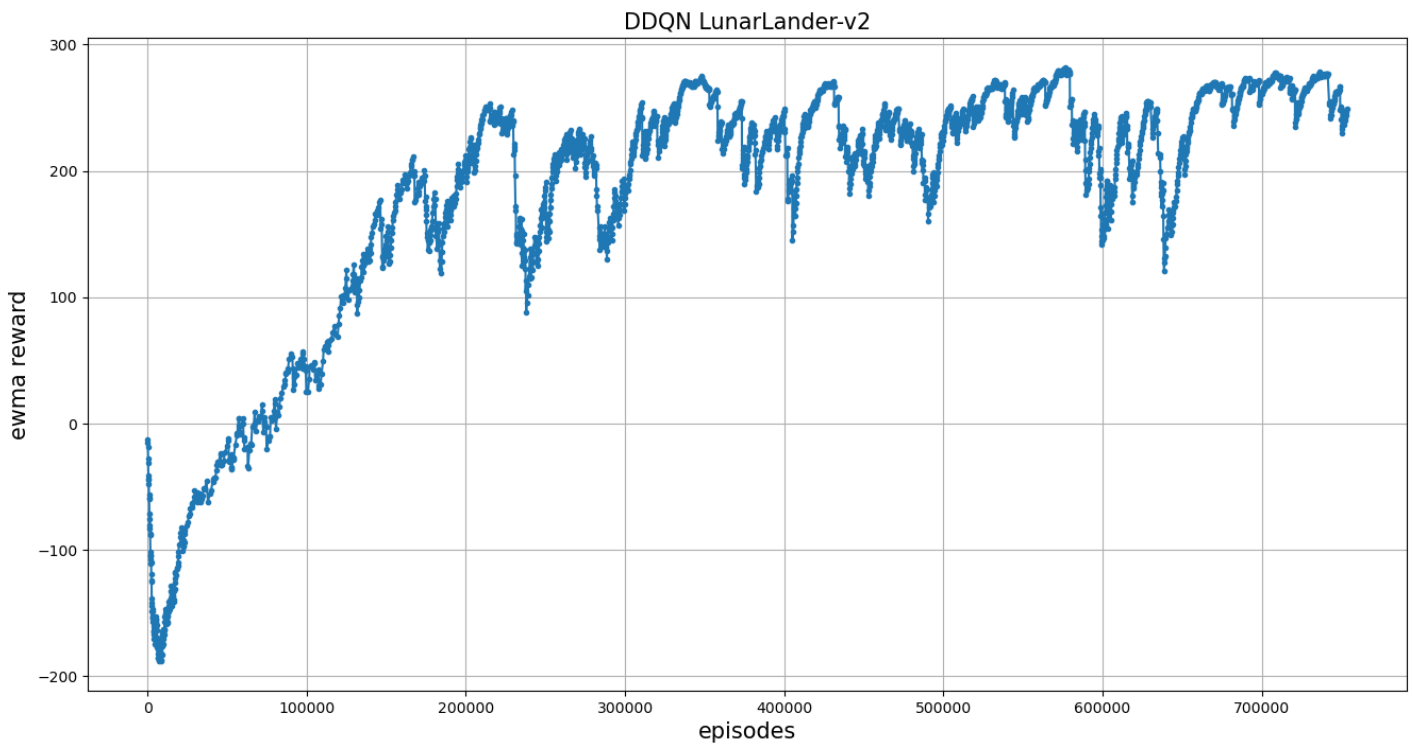
        # bp
        self._optimizer.zero_grad()
        loss.backward()
        nn.utils.clip_grad_norm_(self._behavior_net.parameters(), 5)
        self._optimizer.step()
```

## DDQN: Average Reward: 250.845

Step: 748660	Episode: 2973	Length: 180	Total reward: 247.85	Ewma reward: 262.27	Epsilon: 0.010
Step: 748867	Episode: 2974	Length: 207	Total reward: 306.52	Ewma reward: 264.48	Epsilon: 0.010
Step: 749059	Episode: 2975	Length: 192	Total reward: 316.95	Ewma reward: 267.11	Epsilon: 0.010
Step: 749232	Episode: 2976	Length: 173	Total reward: 261.13	Ewma reward: 266.81	Epsilon: 0.010
Step: 749330	Episode: 2977	Length: 98	Total reward: -102.56	Ewma reward: 248.34	Epsilon: 0.010
Step: 749495	Episode: 2978	Length: 165	Total reward: 255.28	Ewma reward: 248.69	Epsilon: 0.010
Step: 749707	Episode: 2979	Length: 212	Total reward: 272.98	Ewma reward: 249.90	Epsilon: 0.010
Step: 749897	Episode: 2980	Length: 190	Total reward: 266.13	Ewma reward: 250.71	Epsilon: 0.010
Step: 750052	Episode: 2981	Length: 155	Total reward: 57.25	Ewma reward: 241.04	Epsilon: 0.010
Step: 750163	Episode: 2982	Length: 111	Total reward: 19.37	Ewma reward: 229.96	Epsilon: 0.010
Step: 750363	Episode: 2983	Length: 200	Total reward: 295.22	Ewma reward: 233.22	Epsilon: 0.010
Step: 750600	Episode: 2984	Length: 237	Total reward: 281.03	Ewma reward: 235.61	Epsilon: 0.010
Step: 750775	Episode: 2985	Length: 175	Total reward: 247.76	Ewma reward: 236.22	Epsilon: 0.010
Step: 750972	Episode: 2986	Length: 197	Total reward: 322.36	Ewma reward: 240.52	Epsilon: 0.010
Step: 751291	Episode: 2987	Length: 319	Total reward: 279.22	Ewma reward: 242.46	Epsilon: 0.010
Step: 751450	Episode: 2988	Length: 159	Total reward: 265.70	Ewma reward: 243.62	Epsilon: 0.010
Step: 751598	Episode: 2989	Length: 148	Total reward: 256.47	Ewma reward: 244.26	Epsilon: 0.010
Step: 751761	Episode: 2990	Length: 163	Total reward: 269.11	Ewma reward: 245.51	Epsilon: 0.010
Step: 751940	Episode: 2991	Length: 179	Total reward: 288.63	Ewma reward: 247.66	Epsilon: 0.010
Step: 752056	Episode: 2992	Length: 116	Total reward: 31.64	Ewma reward: 236.86	Epsilon: 0.010
Step: 752293	Episode: 2993	Length: 237	Total reward: 306.00	Ewma reward: 240.32	Epsilon: 0.010
Step: 752507	Episode: 2994	Length: 214	Total reward: 261.92	Ewma reward: 241.40	Epsilon: 0.010
Step: 752680	Episode: 2995	Length: 173	Total reward: 275.34	Ewma reward: 243.09	Epsilon: 0.010
Step: 752861	Episode: 2996	Length: 181	Total reward: 269.91	Ewma reward: 244.44	Epsilon: 0.010
Step: 753077	Episode: 2997	Length: 216	Total reward: 282.25	Ewma reward: 246.33	Epsilon: 0.010
Step: 753356	Episode: 2998	Length: 279	Total reward: 292.02	Ewma reward: 248.61	Epsilon: 0.010
Step: 753503	Episode: 2999	Length: 147	Total reward: 265.00	Ewma reward: 249.43	Epsilon: 0.010

Start Testing  
total reward: 245.37  
total reward: 270.37  
total reward: 278.51  
total reward: 262.98  
total reward: 29.58  
total reward: 271.04  
total reward: 282.04  
total reward: 280.66  
total reward: 292.10  
total reward: 295.78  
Average Reward 250.84409525944935  
(pytorch) jackkuo@lab708-Default-string:~/DLPhw6\$

## Training ewma reward/episode of DDQN:





### 13. Performance

#### DQN: Average Reward: 289.75

```
Step: 532439 Episode: 1977 Length: 229 Total reward: 305.00 Ewma reward: 273.35 Epsilon: 0.010
Step: 532641 Episode: 1978 Length: 202 Total reward: 271.09 Ewma reward: 273.24 Epsilon: 0.010
Step: 532816 Episode: 1979 Length: 175 Total reward: 236.59 Ewma reward: 271.41 Epsilon: 0.010
Step: 533000 Episode: 1980 Length: 184 Total reward: 248.19 Ewma reward: 270.25 Epsilon: 0.010
Step: 533209 Episode: 1981 Length: 209 Total reward: 254.25 Ewma reward: 269.45 Epsilon: 0.010
Step: 533409 Episode: 1982 Length: 200 Total reward: 271.11 Ewma reward: 269.53 Epsilon: 0.010
Step: 533604 Episode: 1983 Length: 195 Total reward: 271.61 Ewma reward: 269.63 Epsilon: 0.010
Step: 533785 Episode: 1984 Length: 181 Total reward: 250.72 Ewma reward: 268.69 Epsilon: 0.010
Step: 533984 Episode: 1985 Length: 199 Total reward: 252.11 Ewma reward: 267.86 Epsilon: 0.010
Step: 534195 Episode: 1986 Length: 211 Total reward: 271.38 Ewma reward: 268.04 Epsilon: 0.010
Step: 534406 Episode: 1987 Length: 211 Total reward: 274.44 Ewma reward: 268.36 Epsilon: 0.010
Step: 534611 Episode: 1988 Length: 205 Total reward: 298.93 Ewma reward: 269.88 Epsilon: 0.010
Step: 534797 Episode: 1989 Length: 186 Total reward: 291.35 Ewma reward: 270.96 Epsilon: 0.010
Step: 535020 Episode: 1990 Length: 223 Total reward: 290.43 Ewma reward: 271.93 Epsilon: 0.010
Step: 535221 Episode: 1991 Length: 201 Total reward: 266.62 Ewma reward: 271.67 Epsilon: 0.010
Step: 535404 Episode: 1992 Length: 183 Total reward: 272.26 Ewma reward: 271.70 Epsilon: 0.010
Step: 535600 Episode: 1993 Length: 196 Total reward: 306.81 Ewma reward: 273.45 Epsilon: 0.010
Step: 535790 Episode: 1994 Length: 190 Total reward: 264.62 Ewma reward: 273.01 Epsilon: 0.010
Step: 536421 Episode: 1995 Length: 631 Total reward: 242.20 Ewma reward: 271.47 Epsilon: 0.010
Step: 536621 Episode: 1996 Length: 200 Total reward: 290.19 Ewma reward: 272.40 Epsilon: 0.010
Step: 536824 Episode: 1997 Length: 203 Total reward: 273.22 Ewma reward: 272.45 Epsilon: 0.010
Step: 537032 Episode: 1998 Length: 208 Total reward: 266.90 Ewma reward: 272.17 Epsilon: 0.010
Step: 537275 Episode: 1999 Length: 243 Total reward: 312.76 Ewma reward: 274.20 Epsilon: 0.010
Start Testing
total reward: 249.31
total reward: 289.95
total reward: 278.94
total reward: 277.21
total reward: 318.31
total reward: 266.27
total reward: 306.73
total reward: 301.82
total reward: 317.72
total reward: 291.27
Average Reward 289.7509202940418
(pytorch) jackkuo@lab708-Default-string:~/DLPhw6$
```

#### DDPG: Average Reward: 282.36

```
Step: 891533 Episode: 2977 Length: 202 Total reward: 282.89 Ewma reward: 273.98
Step: 891713 Episode: 2978 Length: 180 Total reward: 248.06 Ewma reward: 272.68
Step: 891886 Episode: 2979 Length: 173 Total reward: 261.66 Ewma reward: 272.13
Step: 892082 Episode: 2980 Length: 196 Total reward: 263.54 Ewma reward: 271.70
Step: 892275 Episode: 2981 Length: 193 Total reward: 279.98 Ewma reward: 272.11
Step: 892469 Episode: 2982 Length: 194 Total reward: 282.47 Ewma reward: 272.63
Step: 892636 Episode: 2983 Length: 167 Total reward: 264.61 Ewma reward: 272.23
Step: 892848 Episode: 2984 Length: 212 Total reward: 291.23 Ewma reward: 273.18
Step: 893045 Episode: 2985 Length: 197 Total reward: 253.63 Ewma reward: 272.20
Step: 893242 Episode: 2986 Length: 197 Total reward: 278.86 Ewma reward: 272.54
Step: 893450 Episode: 2987 Length: 208 Total reward: 271.23 Ewma reward: 272.47
Step: 893870 Episode: 2988 Length: 420 Total reward: 289.62 Ewma reward: 273.33
Step: 894073 Episode: 2989 Length: 203 Total reward: 265.83 Ewma reward: 272.95
Step: 894278 Episode: 2990 Length: 205 Total reward: 299.00 Ewma reward: 274.26
Step: 894508 Episode: 2991 Length: 230 Total reward: 293.64 Ewma reward: 275.22
Step: 894675 Episode: 2992 Length: 167 Total reward: 268.60 Ewma reward: 274.89
Step: 894870 Episode: 2993 Length: 195 Total reward: 274.07 Ewma reward: 274.85
Step: 895077 Episode: 2994 Length: 207 Total reward: 304.46 Ewma reward: 276.33
Step: 895328 Episode: 2995 Length: 251 Total reward: 292.37 Ewma reward: 277.13
Step: 895500 Episode: 2996 Length: 172 Total reward: 285.88 Ewma reward: 277.57
Step: 895706 Episode: 2997 Length: 206 Total reward: 277.85 Ewma reward: 277.59
Step: 895931 Episode: 2998 Length: 225 Total reward: 277.64 Ewma reward: 277.59
Step: 896082 Episode: 2999 Length: 151 Total reward: 270.53 Ewma reward: 277.24
Start Testing
total reward: 252.81
total reward: 285.40
total reward: 281.74
total reward: 280.37
total reward: 307.27
total reward: 268.19
total reward: 283.82
total reward: 295.76
total reward: 290.24
total reward: 278.04
Average Reward 282.3637702581553
(pytorch) jackkuo@lab708-Default-string:~/DLPhw6$
```